

Prof. Wei
Prof. Boddeti

TODAY

Value Iteration

Policy Iteration

- Policy Evaluation
- Policy Extraction

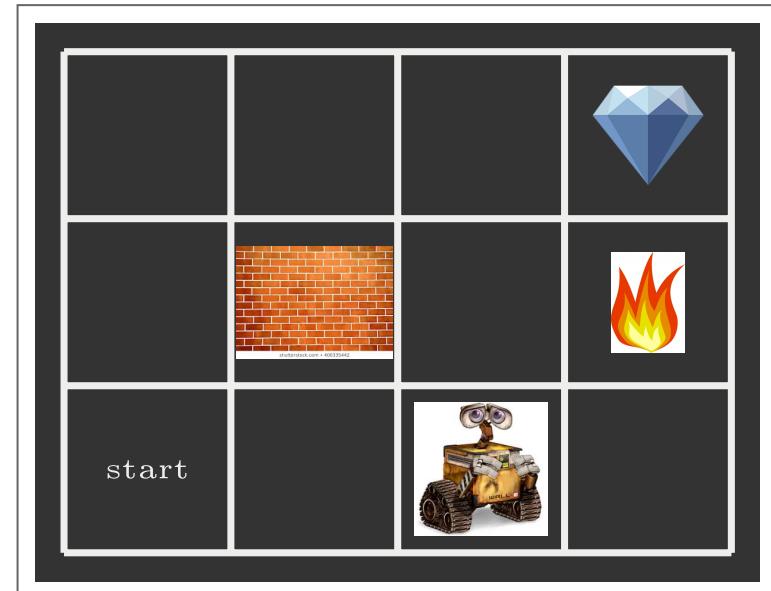
EXAMPLE: GRID WORLD

A maze like problem

- The agent lives in a grid
- Walls block the agent's path
- Noisy movement: actions do not always go as planned
- 80% of the time, the action North takes the agent North (if there is no wall there)
- 10% of the time, North takes the agent West; 10% East
- If there is a wall in the direction the agent would have been taken, the agent stays put

The agent receives rewards each time step

- Small "living" reward each step (can be negative)
- Big rewards come at the end (good or bad)

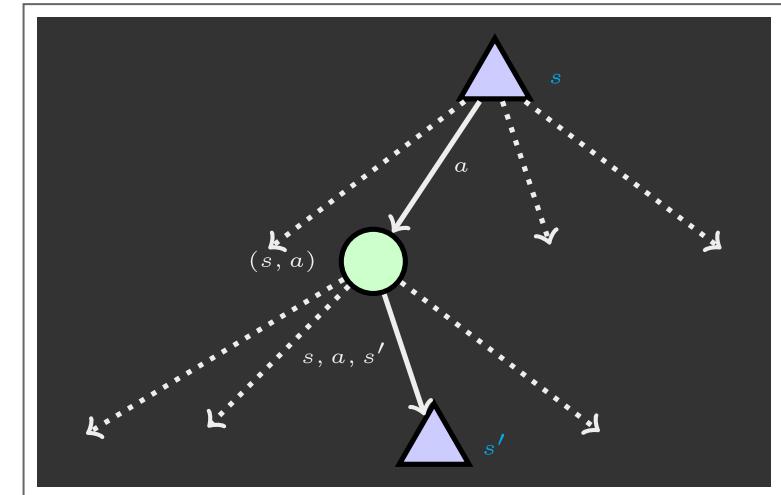


Goal: maximize sum of rewards

RECAP: DEFINING MDPs

Markov decision processes:

- Set of states S
- Start state s_0
- Set of actions A
- Transitions $P(s'|s, a)$ (or $T(s, a, s')$)
- Rewards $R(s, a, s')$ (and discount γ)



MDP quantities so far:

- Policy = Choice of action for each state
- Utility = sum of (discounted) rewards

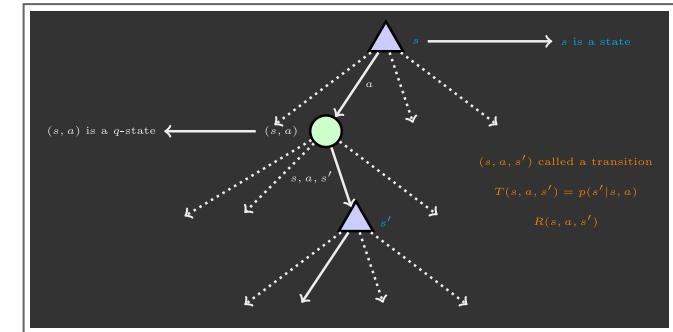
SOLVING MDPS

OPTIMAL QUANTITIES

The value (utility) of a state s : $V^*(s)$ = expected utility starting in s and acting optimally

The value (utility) of a q -state (s, a) : $Q^*(s, a)$ = expected utility starting out having taken action a from state s and (thereafter) acting optimally

The optimal policy: $\pi^*(s)$ = optimal action from state s

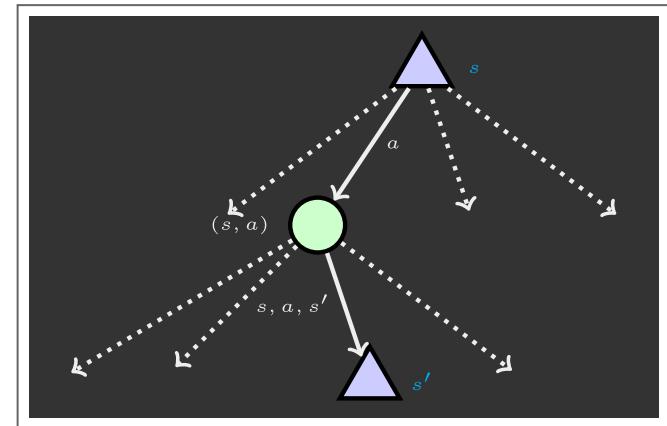


VALUES OF STATES

Fundamental operation: compute the (expectimax) value of a state

- Expected utility under optimal action
- Average sum of (discounted) rewards
- This is just what expectimax computed

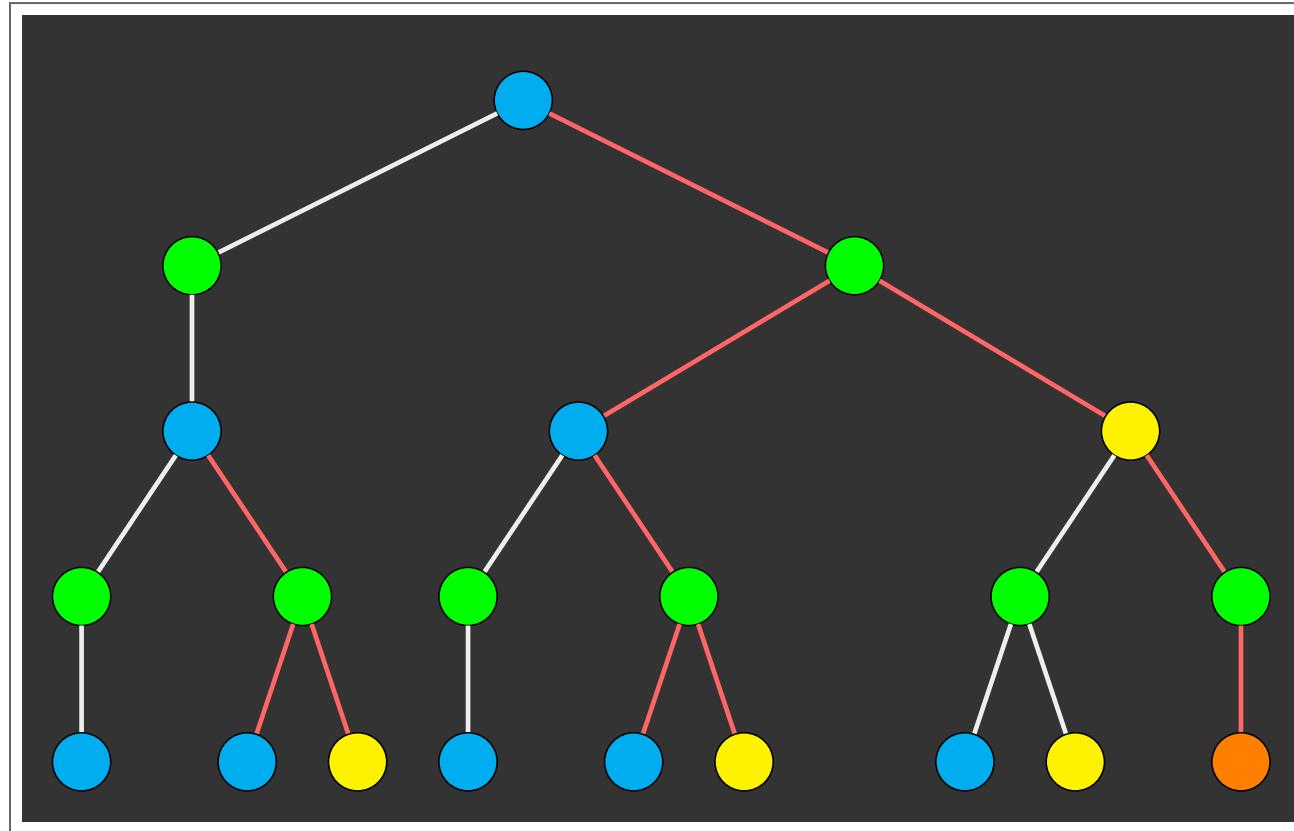
Recursive definition of value:



$$V^*(s) = \max_a Q^*(s, a) \quad (1)$$

$$Q^*(s, a) = \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V^*(s')] \quad (2)$$

RACING SEARCH TREE



RACING SEARCH TREE

We are doing way too much work with expectimax.

Problem 1: States are repeated

- Idea: Only compute needed quantities once

Problem 2: Tree goes on forever

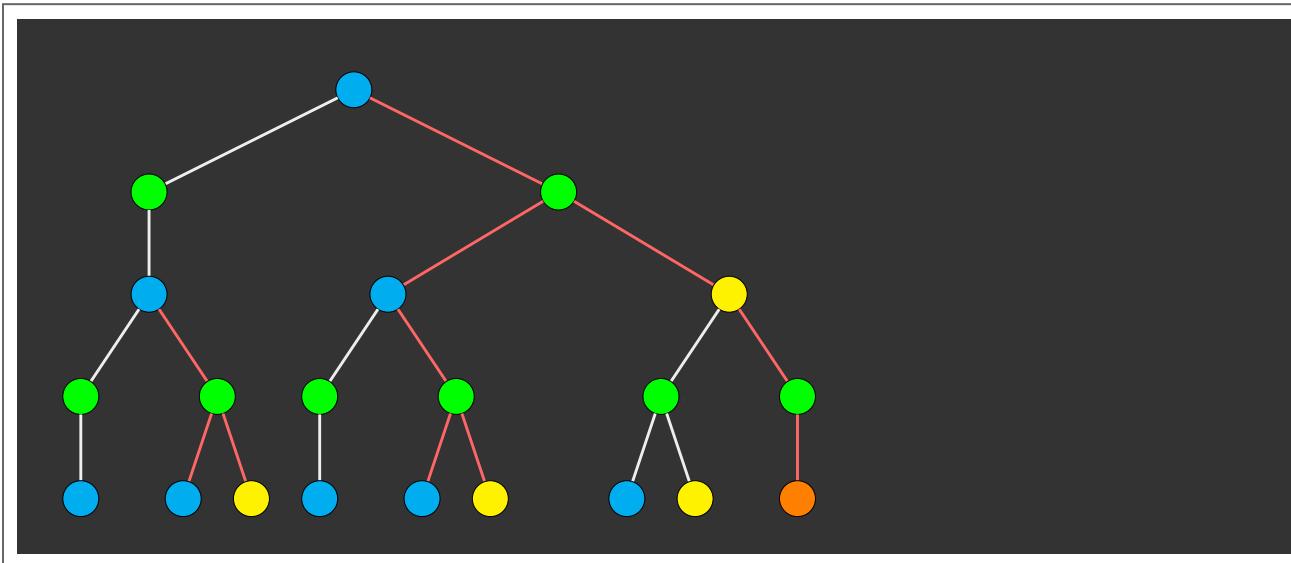
- Idea: Do a depth-limited computation, but with increasing depths until change is small
- Note: deep parts of the tree eventually do not matter if $\gamma < 1$

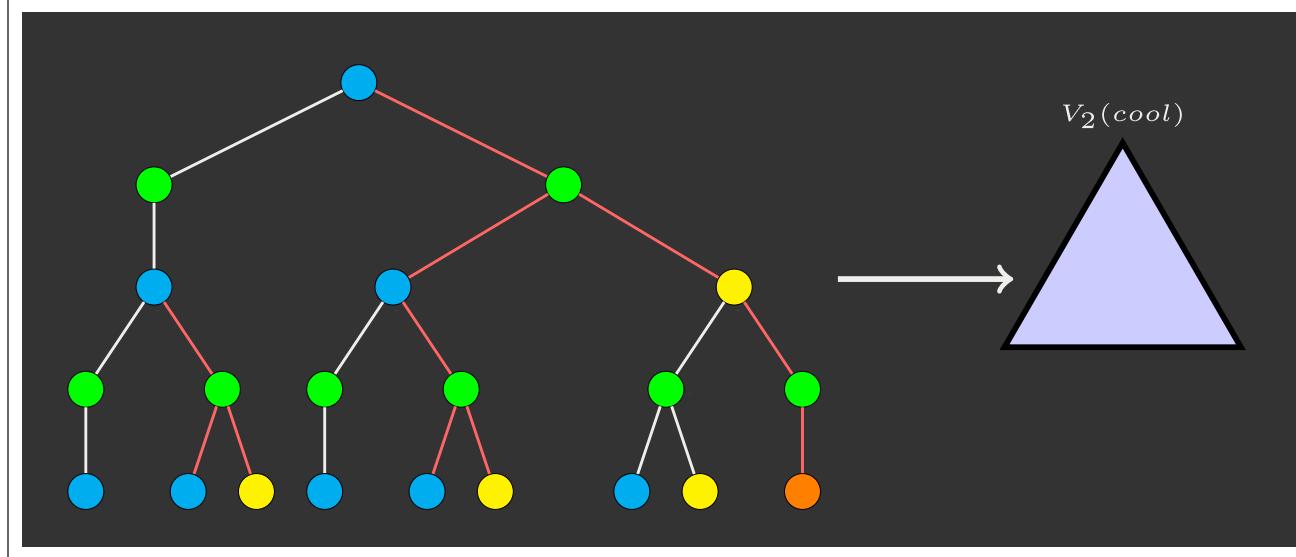
TIME-LIMITED VALUES

Key idea: time-limited values

Define $V_k(s)$ to be the optimal value of s if the game ends in k more time steps

- Equivalently, it is what a depth- k expectimax would give from s





VALUE ITERATION

VALUE ITERATION

Start with $V_0(s) = 0$: no time steps left means an expected reward sum of zero

Given vector of $V_k(s)$ values, do one ply of expectimax from each state:

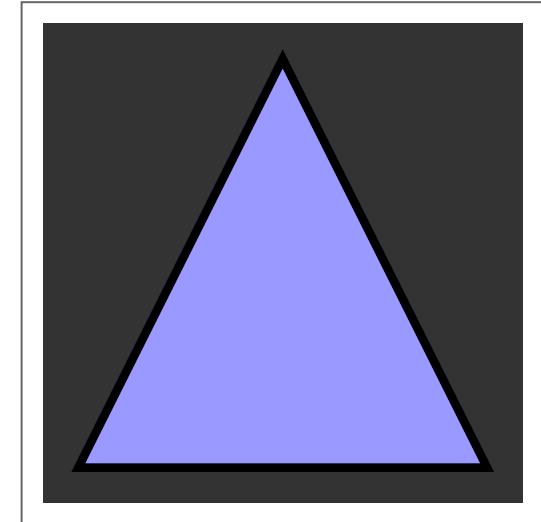
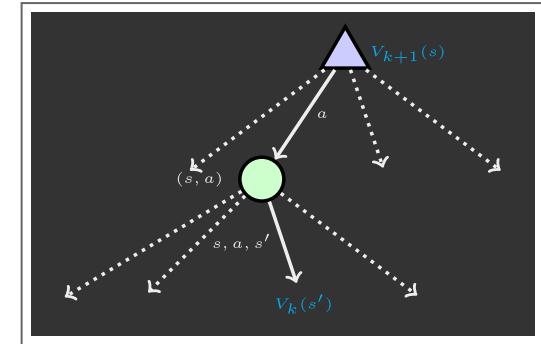
$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k(s')]$$

Repeat until convergence

Complexity of each iteration: $\mathcal{O}(S^2 A)$

Theorem: will converge to unique optimal values

- Basic idea: approximations get refined towards optimal values
- Policy may converge long before values do



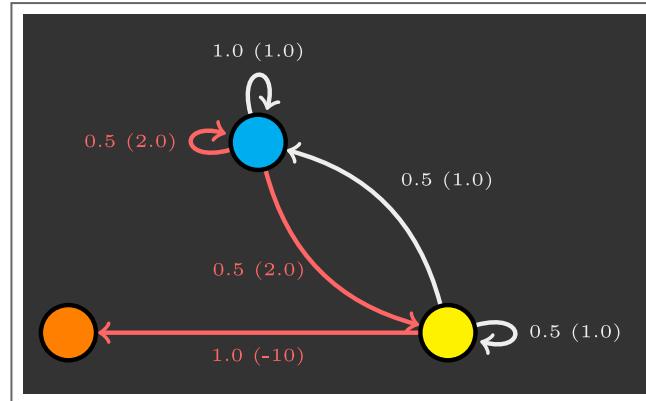
EXAMPLE VALUE ITERATION

Cool Warm Overheated

V_2	3.5	2.5	0
-------	-----	-----	---

V_1	2	1	0
-------	---	---	---

V_0	0	0	0
-------	---	---	---



$$V_{k+1}(s) \leftarrow \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + V_k(s')]$$

Assume no discount

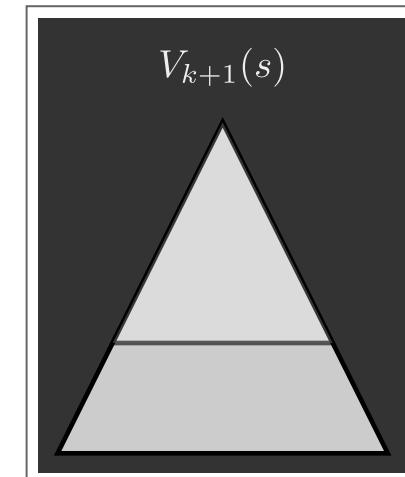
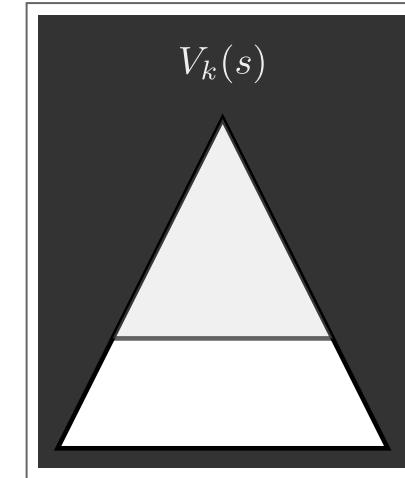
CONVERGENCE

How do we know the V_k vectors are going to converge?

Case 1: If the tree has maximum depth M , then V_M holds the actual untruncated values

Case 2: If the discount is less than 1

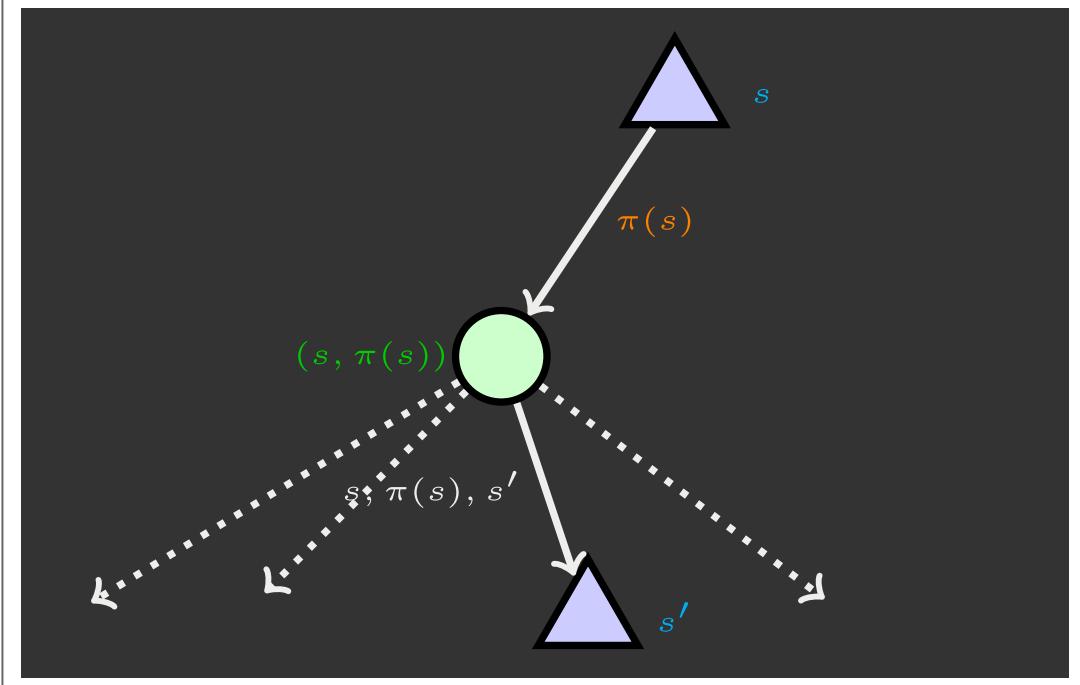
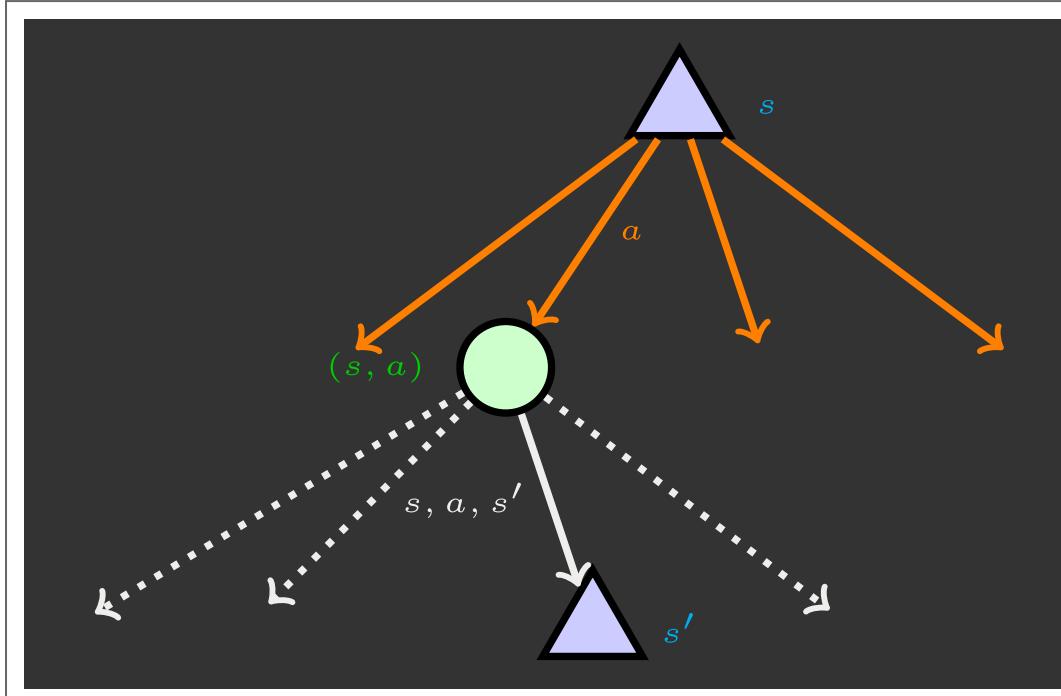
- Sketch: For any state V_k and V_{k+1} can be viewed as depth $k + 1$ expectimax results in nearly identical search trees
- The difference is that on the bottom layer, V_{k+1} has actual rewards while V_k has zeros
- That last layer is at best all R_{MAX}
- It is at worst R_{MIN}
- But everything is discounted by γ^k that far out
- So V_k and V_{k+1} are at most $\gamma^k \max |R|$ different



- So as k increases, the values converge

POLICY EVALUATION

FIXED POLICIES



Expectimax trees max over all actions to compute the optimal values

If we fixed some policy $\pi(s)$, then the tree would be simpler – only one action per state

- ... though the tree's value would depend on which policy we fixed

UTILITIES FOR A FIXED POLICY

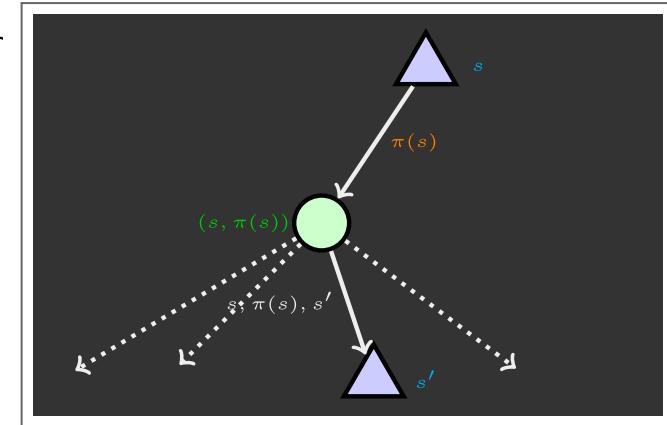
Another basic operation: compute the utility of a state s under a fixed (generally non-optimal) policy

Define the utility of a state s , under a fixed policy π :

$V^\pi(s)$ = expected total discounted rewards starting in s and following π

Recursive relation (one-step look-ahead / Bellman equation):

$$V^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V^\pi(s')]$$



POLICY EVALUATION

How do we calculate the V 's for a fixed policy π ?

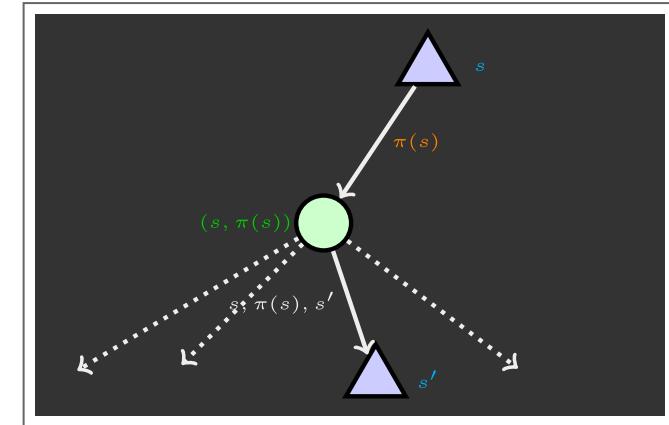
Idea 1: Turn recursive Bellman equations into updates (like value iteration)

$$V_0^\pi(s) = 0$$
$$V_{k+1}^\pi(s) = \sum_{s'} T(s, \pi(s), s') [R(s, \pi(s), s') + \gamma V_k^\pi(s')]$$

Efficiency: $\mathcal{O}(S^2)$ per iteration

Idea 2: Without the maxes, the Bellman equations are just a linear system

- Solve with Matlab (or your favorite linear system solver)



POLICY EXTRACTION

COMPUTING VALUES FROM ACTIONS

Imagine we have the optimal values $V^*(s)$

How should we act?

- It is not obvious

We need to do a mini-expectimax (one step)

$$\pi^*(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s')[R(s, a, s') + \gamma V^*(s')]$$

This is called [policy extraction](#), since it gets the policy implied by the values.

COMPUTING ACTIONS FROM Q-VALUES

Imagine we have the optimal q-values: $Q^*(s, a)$

How should we act?

- Completely trivial to decide !!

$$\pi^*(s) = \operatorname{argmax}_a Q^*(s, a)$$

Important lesson: actions are easier to select from q-values than values !!

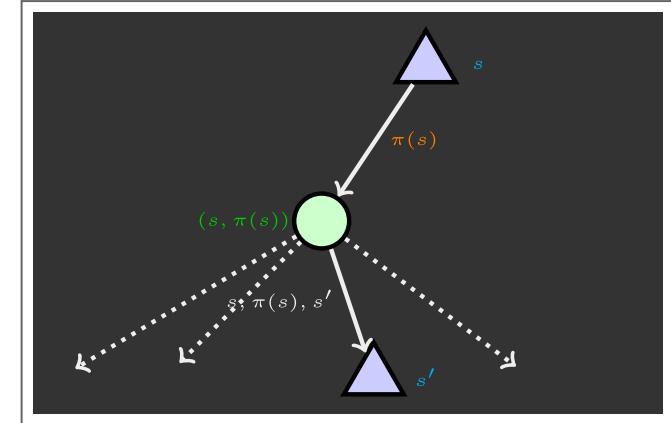
POLICY ITERATION

PROBLEMS WITH VALUE ITERATION

Value iteration repeats the Bellman updates:

$$V_{k+1}^* = \max_a \sum_{s'} T(s, a, s') [R(s, a, s') + \gamma V_k^*(s')]$$

Problem 1: It is slow – $\mathcal{O}(S^2 A)$ per iteration



Problem 2: The "max" at each state rarely changes

Problem 3: The policy often converges long before the values

POLICY ITERATION

Alternative approach for optimal values:

- **Step 1: Policy evaluation:** calculate utilities for some fixed policy (not optimal utilities) until convergence
- **Step 2: Policy improvement:** update policy using one-step look-ahead with resulting converged (but not optimal) utilities as future values
- Repeat steps until policy converges

This is policy iteration

- It is still optimal !!
- Can converge (much) faster under some conditions

POLICY ITERATION

Evaluation: For fixed current policy π , find values with policy evaluation:

- Iterate until values converge:

$$V_{k+1}^{\pi_i}(s) = \sum_{s'} T(s, \pi_i(s), s')[R(s, \pi_i(s), s') + \gamma V_k^{\pi_i}(s')]$$

Improvement: For fixed values, get a better policy using policy extraction

- One-step look-ahead:

$$\pi_{i+1}(s) = \operatorname{argmax}_a \sum_{s'} T(s, a, s')[R(s, a, s') + \gamma V^{\pi_i}(s')]$$

COMPARISON

Both value iteration and policy iteration compute the same thing (all optimal values)

In value iteration:

- Every iteration updates both the values and (implicitly) the policy
- We do not track the policy, but taking the max over actions implicitly recomputes it.

In policy iteration:

- We do several passes that update utilities with fixed policy (each pass is fast because we consider only one action, not all of them)
- After the policy is evaluated, a new policy is chosen (slow like a value iteration pass)
- The new policy will be better (or we are done)

Both are dynamic programs for solving MDPs

SUMMARY: MDP ALGORITHMS

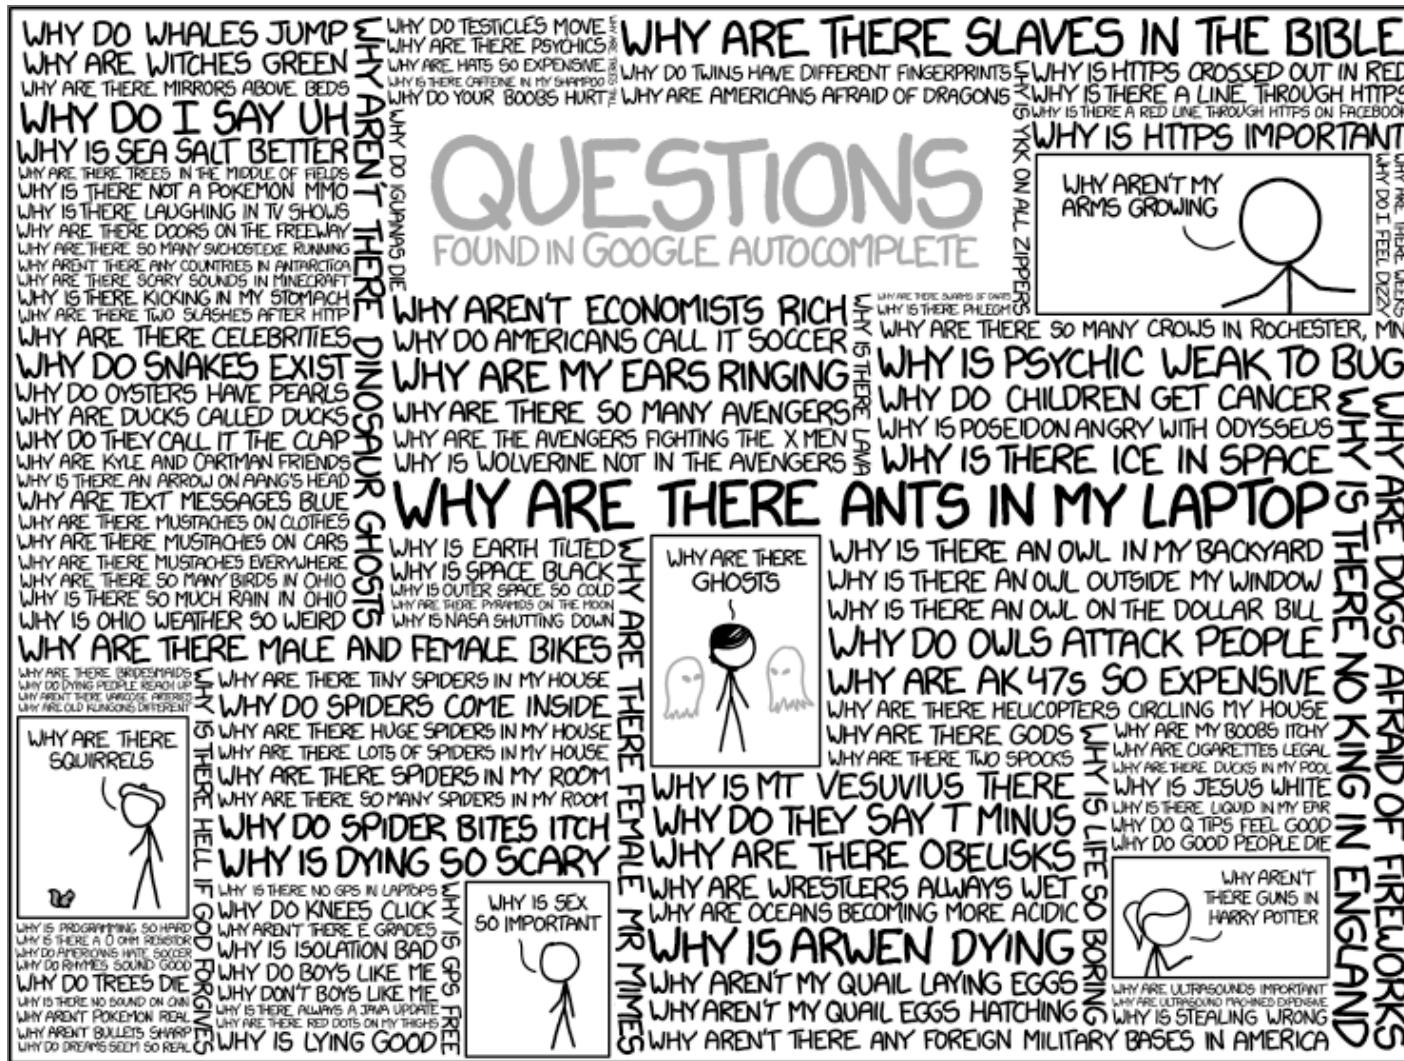
So you want to ...

- Compute optimal values: use value iteration or policy iteration
- Compute values for a particular policy: use policy evaluation
- Turn your values into a policy: use policy extraction (one-step lookahead)

These all look the same ...

- They basically are – they are all variations of Bellman updates
- They all use one-step lookahead expectimax fragments
- They differ only in whether we plug in a fixed policy or max over actions

Q & A



XKCD

Speaker notes

